Assignment 1: Write a SELECT query to retrieve all columns from a 'customers' table, and modify it to return only the customer name and email address for customers in a specific city.

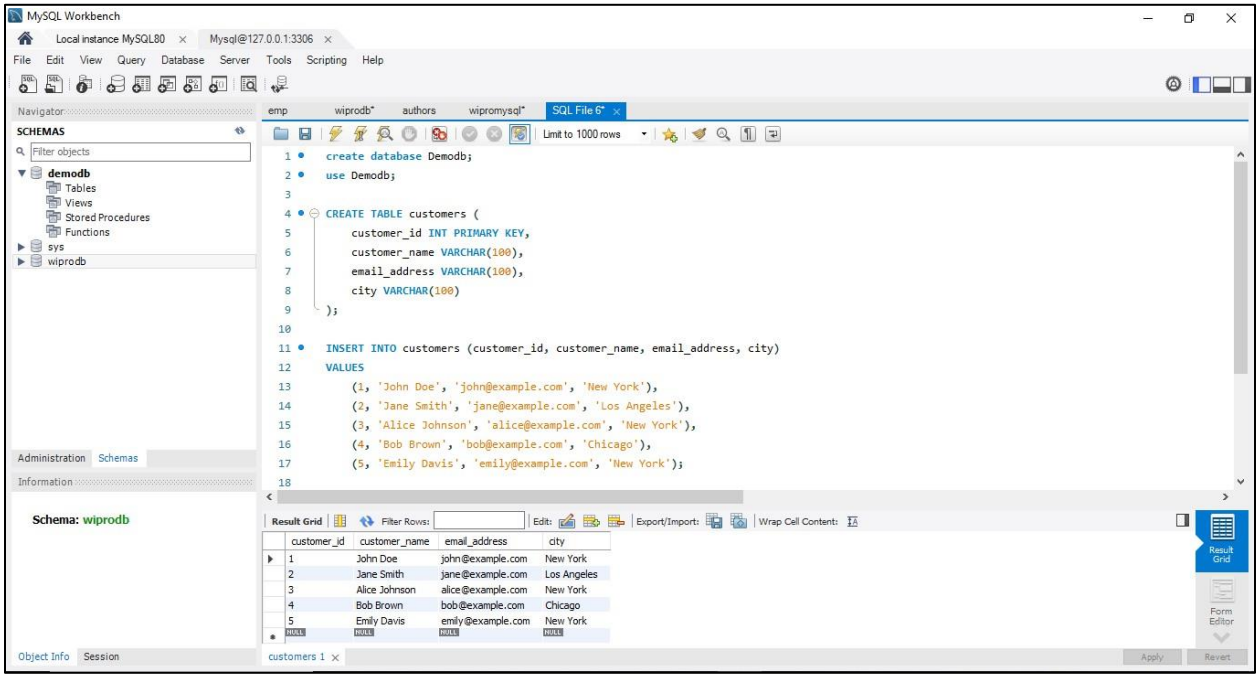

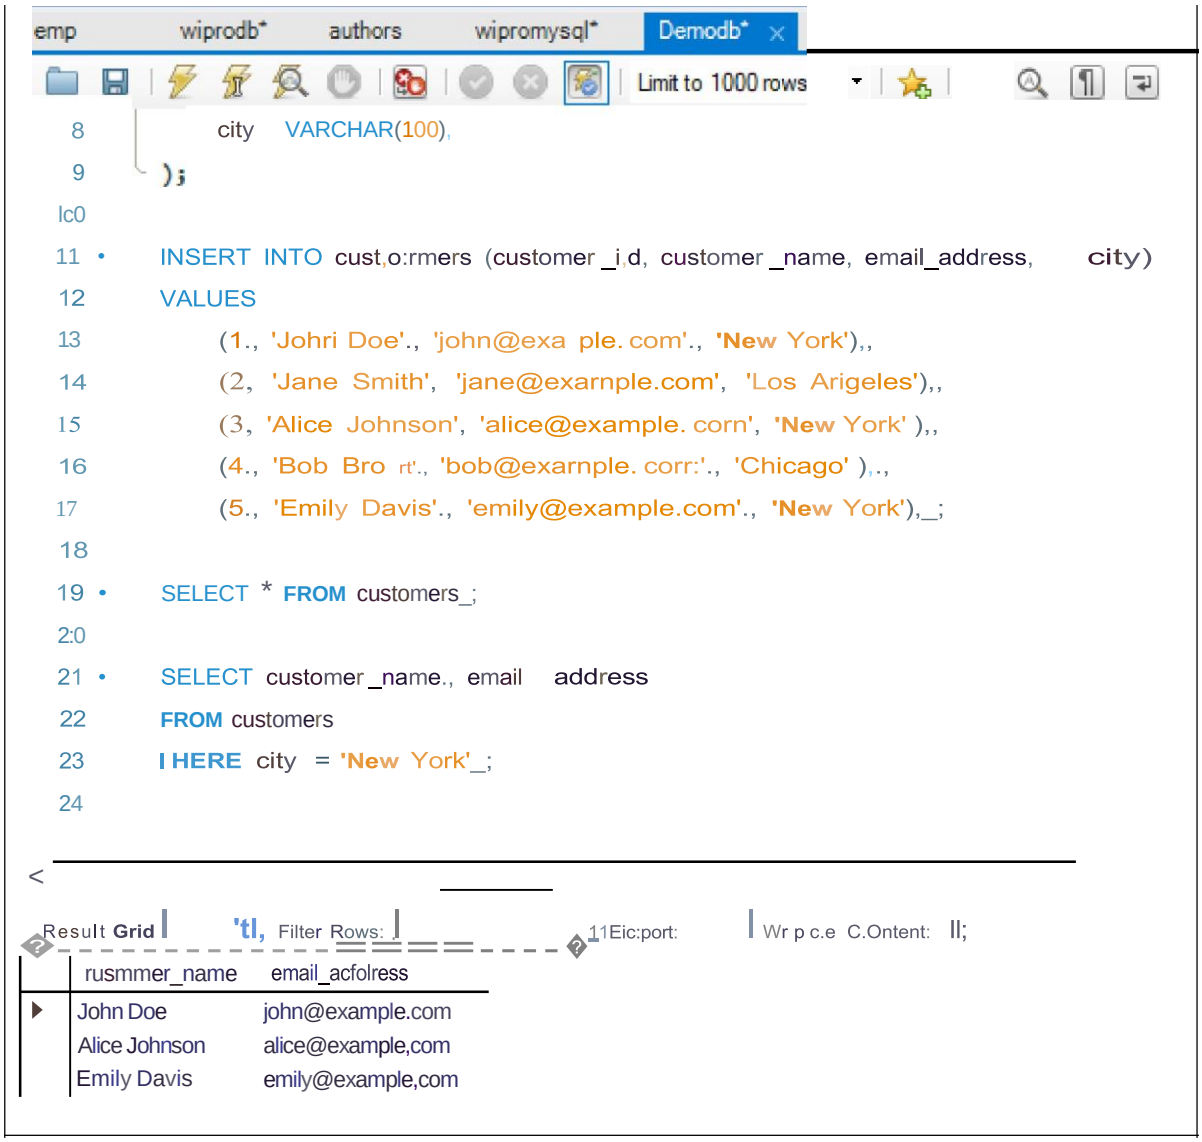
emp
wiprodb*
authors
wipromysql*
Demodb*
Limit to 1000 rows
city VARCHAR(100),
);
INSERT INTO cust,o:rmers (customer _i,d, customer _name, email_address, city)
VALUES
(1., 'Johri Doe'., 'john@exa ple. com'., 'New York'),,
(2, 'Jane Smith', 'jane@exarnple.com', 'Los Arigeles'),,
(3, 'Alice Johnson', 'alice@example. corn', 'New York'),,
(4., 'Bob Bro rt'., 'bob@exarnple. corr:'., 'Chicago'),.,
(5., 'Emily Davis'., 'emily@example.com'., 'New York'),_;
SELECT * FROM customers_;
SELECT customer _name., email address
FROM customers
I HERE city = 'New York'_;
Result Grid
'tl, Filter Rows:
11Eic:port:
Wr p c.e C.Ontent:
rusmmer_name
email_acfolress
John Doe
john@example.com
Alice Johnson
alice@example,com
Emily Davis
emily@example,com

Assignment 2: Craft a query using an INNER JOIN to combine 'orders' and 'customers' tables for customers in a specified region, and a LEFT JOIN to display all customers including those without orders.

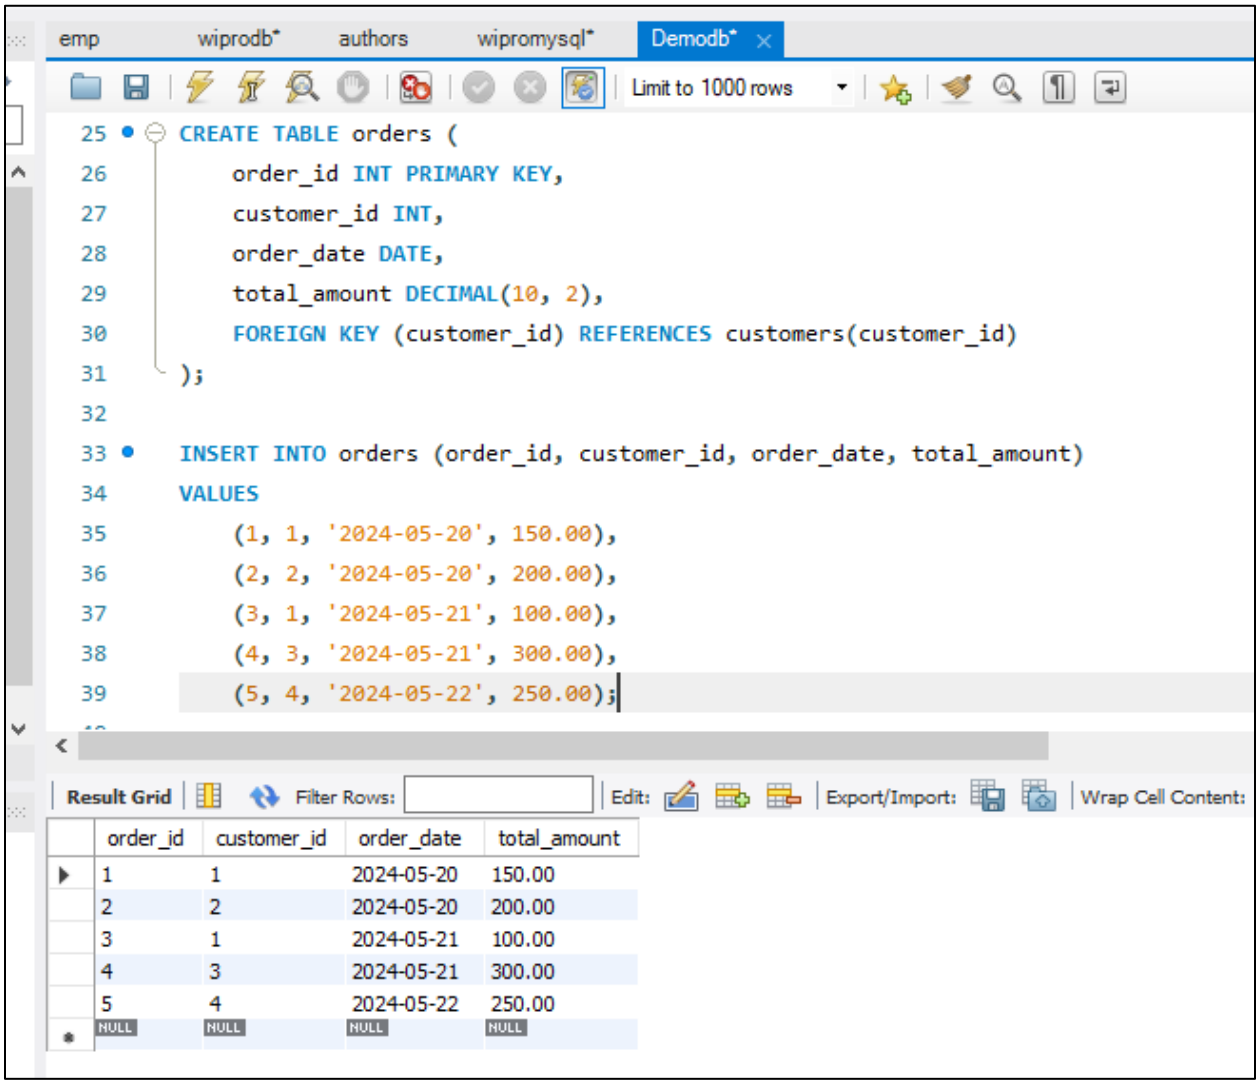

```
40
41 •        select "' from or,dersj
42 •        select"' from customers_;
43
44 •    SELECT c.'", o.orsder_id, o.orsder ,date
45     FROM custo ers c
46     LEFT JOIN or,ders o ON c. customer i\d    o.customer i,d
47     I HERE c. city = 'city'_;
48
```

<

llesul[t **Grid** ı D Filter Rows: ı E>!port: ı Wr p **c.e**l C.Onieni:

rnsromer_id	rnsromer_name	email_address	city	order_id	order_date

Result 9' x

Output

oil | Action Outpl!Il H

	Time	Ac:IJio:ri		
O	f;1 14:20:50	select• from cl!lslomers LIMIT 0, 1000	5 row(s)retl!lmed	
O	68 14:21:03	SELECT c.·, o.order_id, o.order_date FROM customersc LEFT JOINorders o ON c.....	0 row(s) returned	

Assignment 3: Utilize a subquery to find customers who have placed orders above the average order value, and write a UNION query to combine two SELECT statements with the same number of columns.

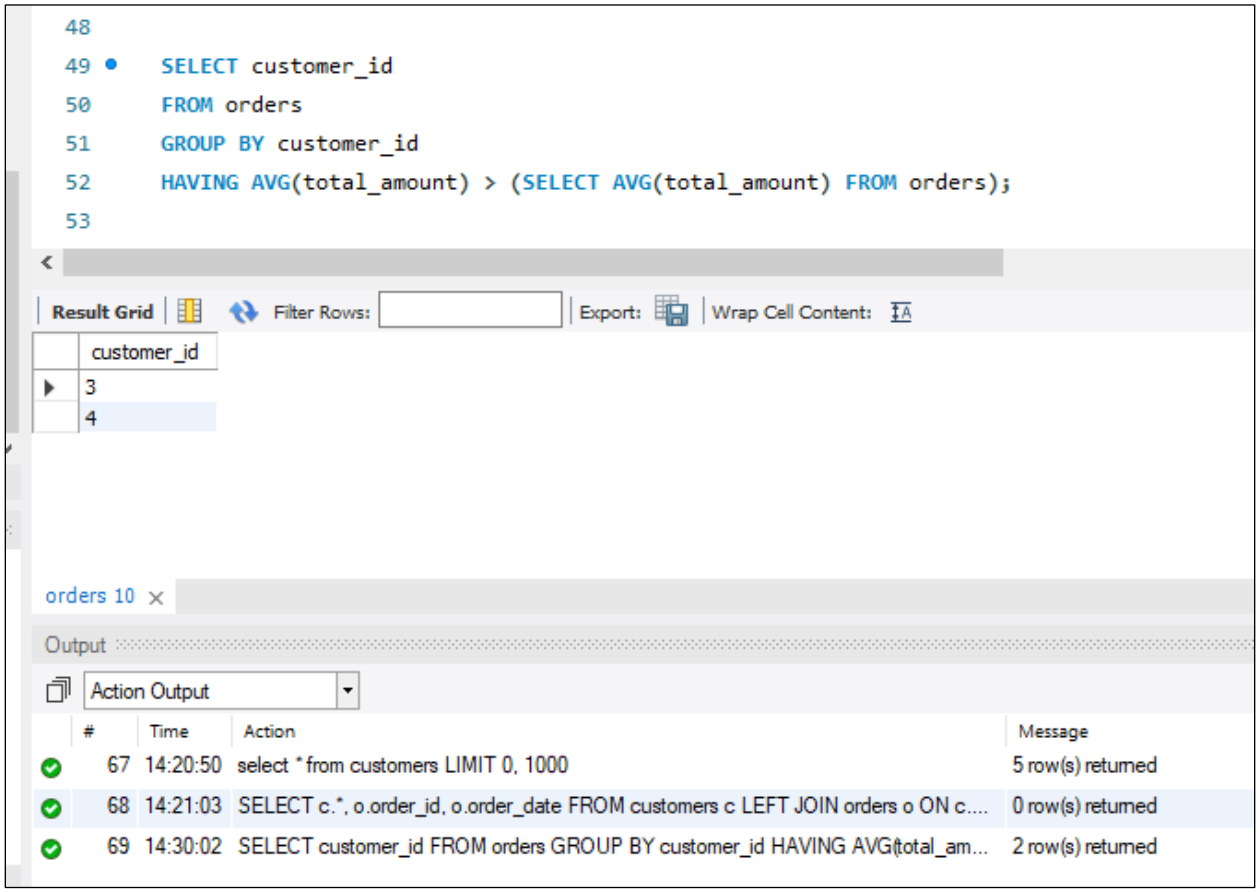

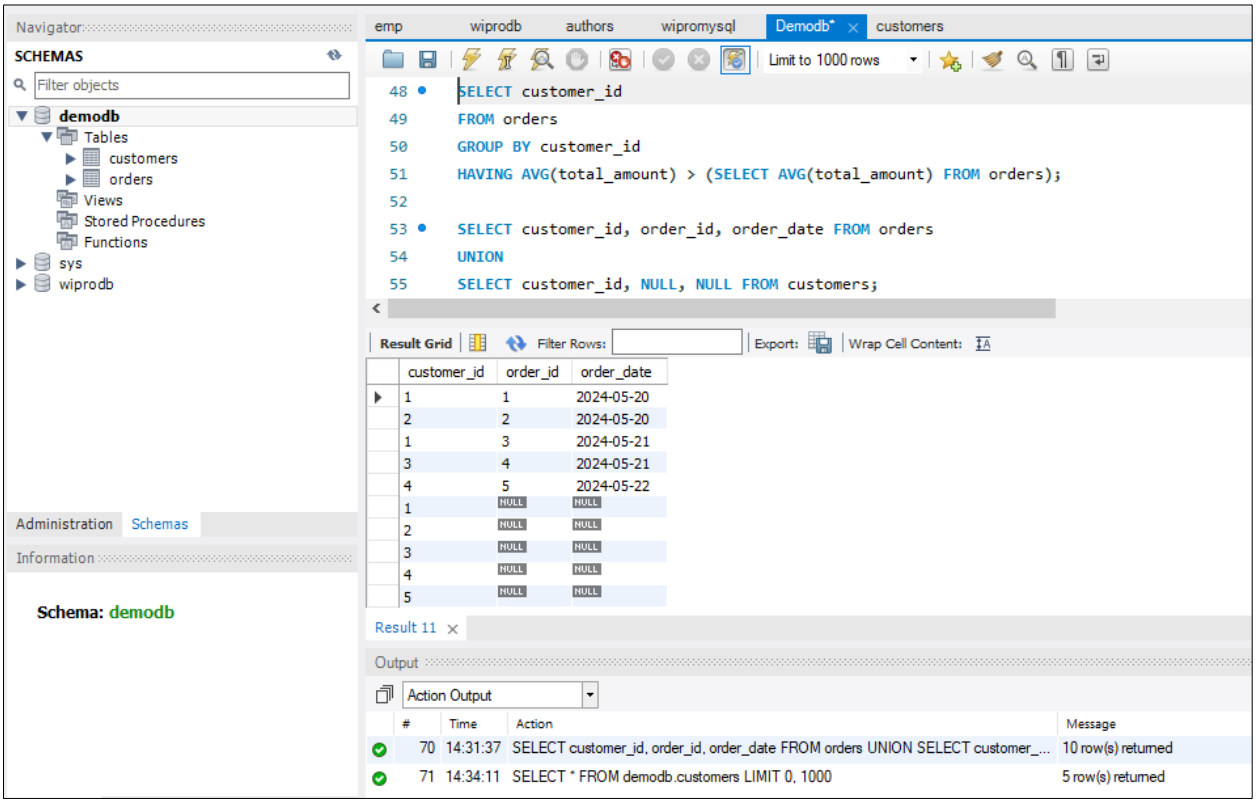

Assignment 4: Compose SQL statements to BEGIN a transaction, INSERT a new record into the 'orders' table, COMMIT the transaction, then UPDATE the 'products' table, and ROLLBACK the transaction.

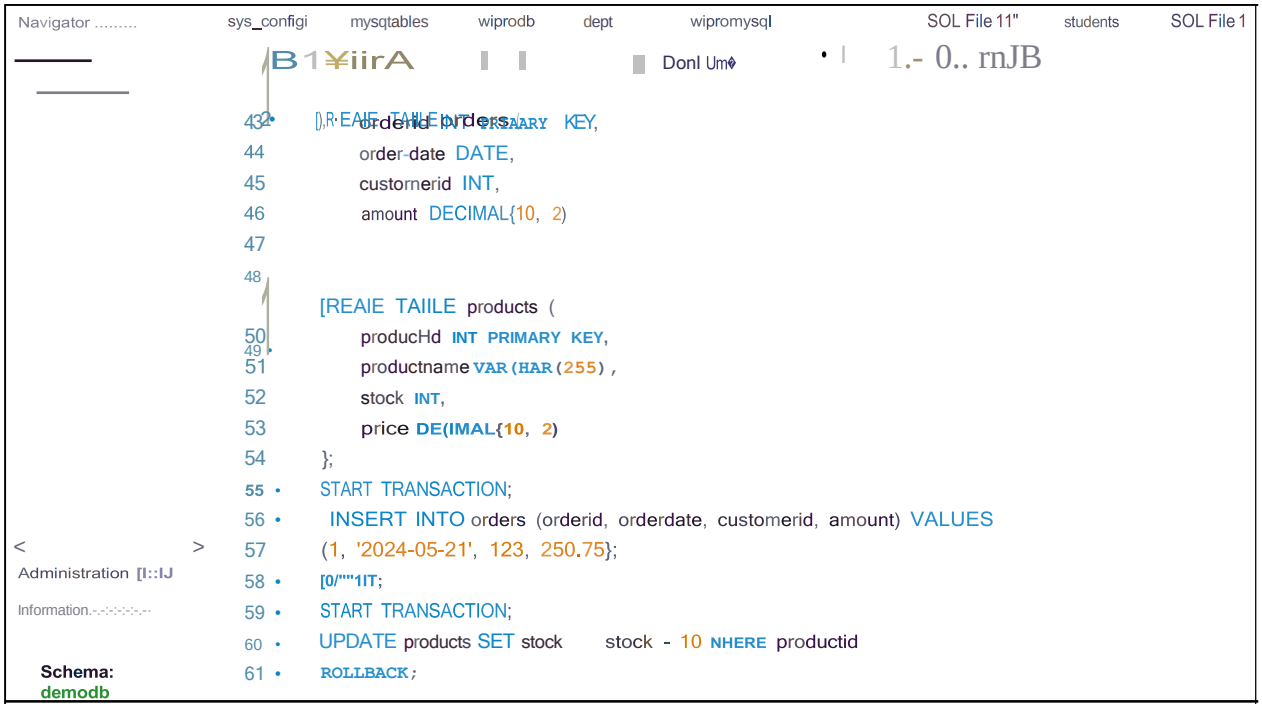
Navigator .....
sys_configi
mysqtables
wiprodb
dept
wipromysql
SOL File 11"
students
SOL File 1
Donl Um
1.- 0.. rnJB
orderid INT PRIMARY KEY,
order-date DATE,
custornerid INT,
amount DECIMAL{10, 2)
[REAIE TAIILE products (
producHd INT PRIMARY KEY,
productname VAR(HAR(255),
stock INT,
price DE(IMAL{10, 2)
};
START TRANSACTION;
INSERT INTO orders (orderid, orderdate, customerid, amount) VALUES
(1, '2024-05-21', 123, 250.75};
[0/""1IT;
START TRANSACTION;
UPDATE products SET stock stock - 10 NHERE productid
ROLLBACK;
Administration [I::IJ
Information.-.-:-:-:-:-.--
Schema:
demodb


SCHEMAS
Filter objects
demodb
librarydb
sys
wiprodb
Tables
Views
Output
ijil AdionOutpu:
Me» i'
1 22:10:44 CREATETABLEorders (ordend INT PRIMARY KEY. orderdate DATE. customend INT. am. Orow(s)affeded
2 22:11:02 CREATE TABLE products (productid INT PRIMARY KEY, pmductname VARCHAR(255), sto... 0 row(s) affected
3 22:11:43 START TRANSACTION Orow(s)affected
4 22:12:45 INSERT INTO o (orderid. orderdate, customerid. amount)VALUES (1. '2(124-05-21'. 123. 250.75) 1 row(s)affected
5 22,UOO COMMIT Orow(,)affeded
6 22:13:45 START TRANSACTION Orow(s) cifeded
7 22:14:02 UPDATEducts SET stock= stock- 70WHERE du::tid = 1 Orow(s)affeded Rows matched: 0 Oianged: 0 Warnings: 0
3 22:14:37 ROLLBACK Orow(,) affected
If5l stored Procedur
Functions

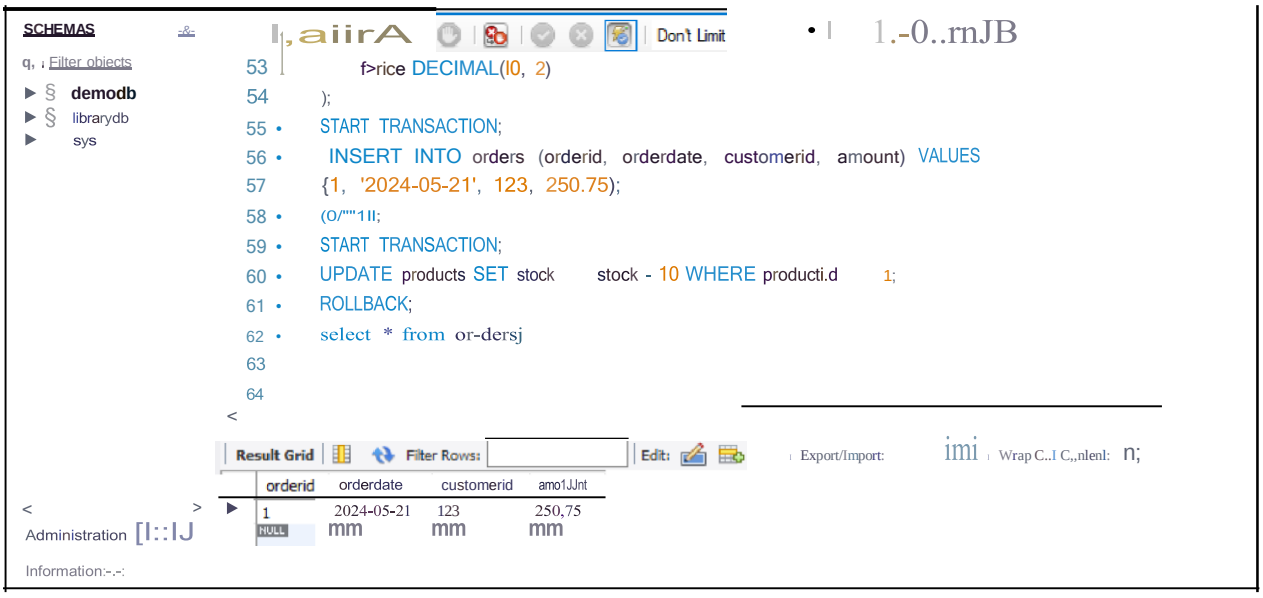
SCHEMAS
Filter objects
demodb
librarydb
sys
Don't Limit
1.-0..rnJB
f>rice DECIMAL(l0, 2)
);
START TRANSACTION;
INSERT INTO orders (orderid, orderdate, customerid, amount) VALUES
{1, '2024-05-21', 123, 250.75);
(0/""1II;
START TRANSACTION;
UPDATE products SET stock stock - 10 WHERE producti.d 1;
ROLLBACK;
select * from or-dersj
Result Grid
Filter Rows:
Edit:
Export/Import:
imi
Wrap C..I C,,nlenl:
n;
orderid
orderdate
customerid
amo1JJnt
1
2024-05-21
123
250,75
NULL
mm
mm
mm
Administration [I::IJ
Information:-.-:

Assignment 5: Begin a transaction, perform a series of INSERTs into 'orders', setting a SAVEPOINT after each, rollback to the second SAVEPOINT, and COMMIT the overall transaction.

Ans:

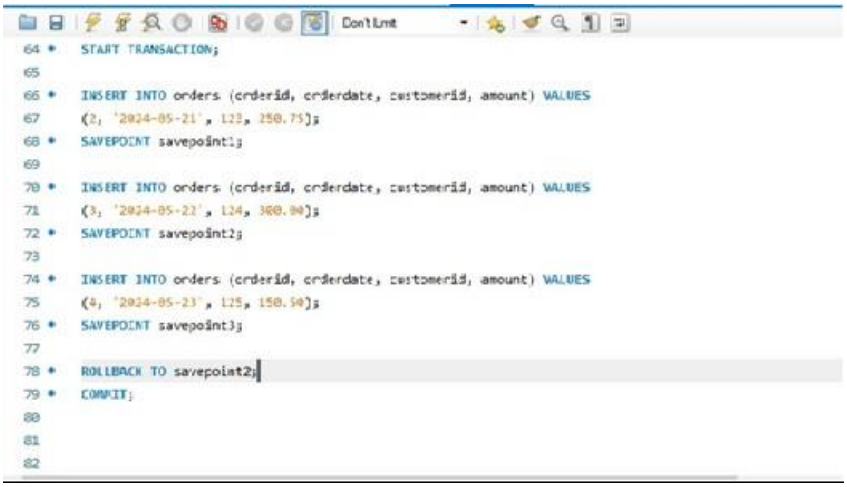

```
START TRANSACTION;

INSERT INTO orders (orderid, orderdate, customerid, amount) VALUES
(2, '2024-05-21', 123, 250.75);
SAVEPOINT savepoint1;

INSERT INTO orders (orderid, orderdate, customerid, amount) VALUES
(3, '2024-05-22', 124, 300.00);
SAVEPOINT savepoint2;

INSERT INTO orders (orderid, orderdate, customerid, amount) VALUES
(4, '2024-05-23', 125, 150.50);
SAVEPOINT savepoint3;

ROLLBACK TO savepoint2;
COMMIT;
```

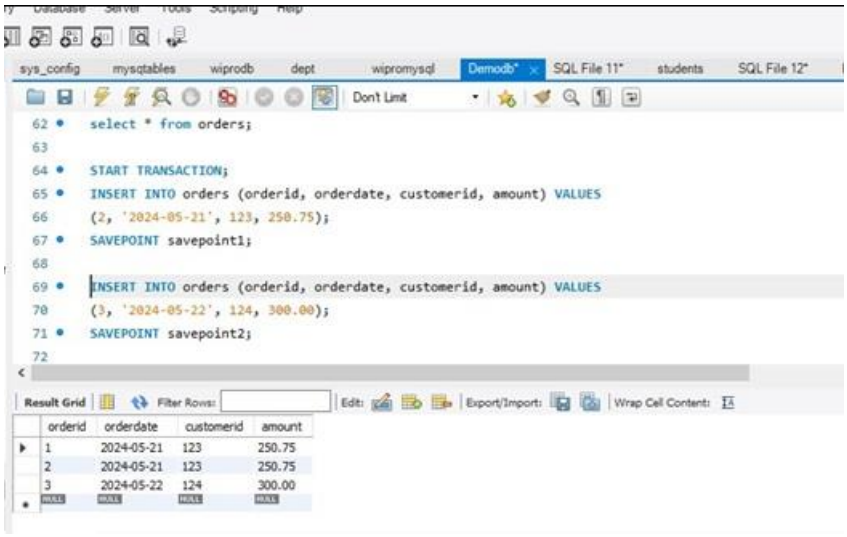

```
select * from orders;

START TRANSACTION;
INSERT INTO orders (orderid, orderdate, customerid, amount) VALUES
(2, '2024-05-21', 123, 250.75);
SAVEPOINT savepoint1;

INSERT INTO orders (orderid, orderdate, customerid, amount) VALUES
(3, '2024-05-22', 124, 300.00);
SAVEPOINT savepoint2;
```

Result Grid

orderid	orderdate	customerid	amount
1	2024-05-21	123	250.75
2	2024-05-21	123	250.75
3	2024-05-22	124	300.00
NULL	NULL	NULL	NULL

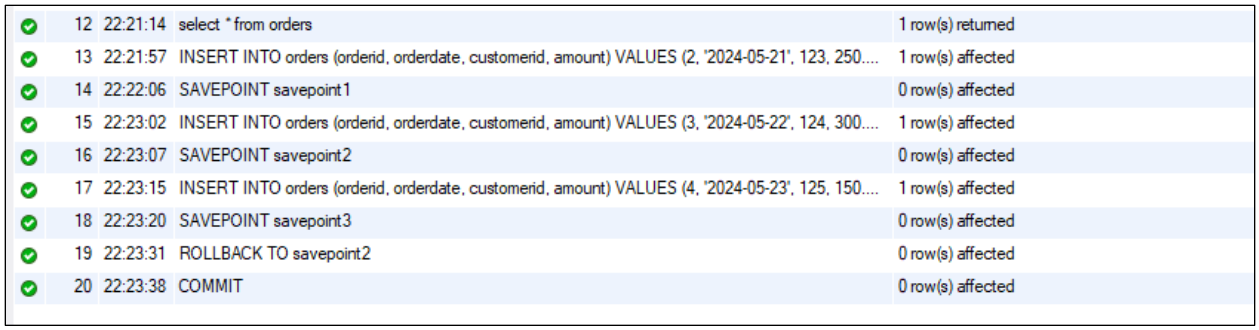

✓	12	22:21:14	select * from orders	1 row(s) returned
✓	13	22:21:57	INSERT INTO orders (orderid, orderdate, customerid, amount) VALUES (2, '2024-05-21', 123, 250....	1 row(s) affected
✓	14	22:22:06	SAVEPOINT savepoint1	0 row(s) affected
✓	15	22:23:02	INSERT INTO orders (orderid, orderdate, customerid, amount) VALUES (3, '2024-05-22', 124, 300....	1 row(s) affected
✓	16	22:23:07	SAVEPOINT savepoint2	0 row(s) affected
✓	17	22:23:15	INSERT INTO orders (orderid, orderdate, customerid, amount) VALUES (4, '2024-05-23', 125, 150....	1 row(s) affected
✓	18	22:23:20	SAVEPOINT savepoint3	0 row(s) affected
✓	19	22:23:31	ROLLBACK TO savepoint2	0 row(s) affected
✓	20	22:23:38	COMMIT	0 row(s) affected

Assignment 6 : Draft a brief report on the use of transaction logs for data recovery and create a hypothetical scenario where a transaction log is instrumental in data recovery after an unexpected shutdown.

Ans:

Report: Leveraging Transaction Logs for Data Recovery in MySQL

Introduction: Transaction logs are vital components of MySQL database management, playing a crucial role in ensuring data integrity and facilitating recovery in the event of system failures or unexpected shutdowns. These logs record every change made to a database, providing a detailed trail of transactions. This report explores the significance of transaction logs for data recovery in MySQL and illustrates their importance through a hypothetical scenario.

The Importance of Transaction Logs for Data Recovery: Transaction logs, specifically the binary logs in MySQL, serve as reliable sources of information for recovering data after a system failure. They maintain a chronological record of all events that modify the database, including inserts, updates, and deletes. By capturing changes before they are permanently written to the database, transaction logs enable the reconstruction of the database to a consistent state, even in the face of unforeseen disruptions.

Key Functions of Transaction Logs in MySQL:

1. **Redo Logging:** MySQL's binary logs capture the changes made to the database, allowing for the replay of transactions that were committed but not yet written to disk at the time of the failure. This process, known as redo logging, ensures that committed transactions are not lost during recovery.

2. **Undo Logging:** InnoDB, the default storage engine for MySQL, uses undo logs to store information necessary to reverse or undo transactions that were in progress but not yet committed at the time of the failure. This capability enables the restoration of the database to its pre-transaction state, maintaining data consistency.

3. **Point-in-Time Recovery:** MySQL's binary logs enable point-in-time recovery, allowing database administrators to restore the database to a specific moment before the failure occurred. By replaying transactions up to the desired timestamp, organizations can minimize data loss and maintain business continuity.

Hypothetical Scenario: Consider a scenario where an e-commerce platform experiences an unexpected shutdown of its MySQL database server due to a hardware failure. As a result, critical customer order data becomes inaccessible, posing a significant operational risk. However, due to the diligent use of MySQL binary logs, the organization can recover the data swiftly and minimize the impact on its operations.

- **Event:** The database server abruptly shuts down, leading to the loss of unsaved changes and potentially jeopardizing the integrity of customer order records.

- **Response:** Upon restarting the database server, database administrators immediately initiate the recovery process using MySQL binary logs. By applying the binary logs, the system reconstructs the database to a consistent state, ensuring that all committed transactions are preserved.

Outcome: Despite the unexpected shutdown, the e-commerce platform successfully restores access to critical customer order data with minimal data loss. The MySQL binary logs prove instrumental in facilitating rapid recovery, demonstrating their indispensable role in ensuring data resilience and business continuity.

Step-by-Step Recovery Process:

1. **Identify the Binary Log Files:** Locate the binary log files generated by MySQL before the crash.

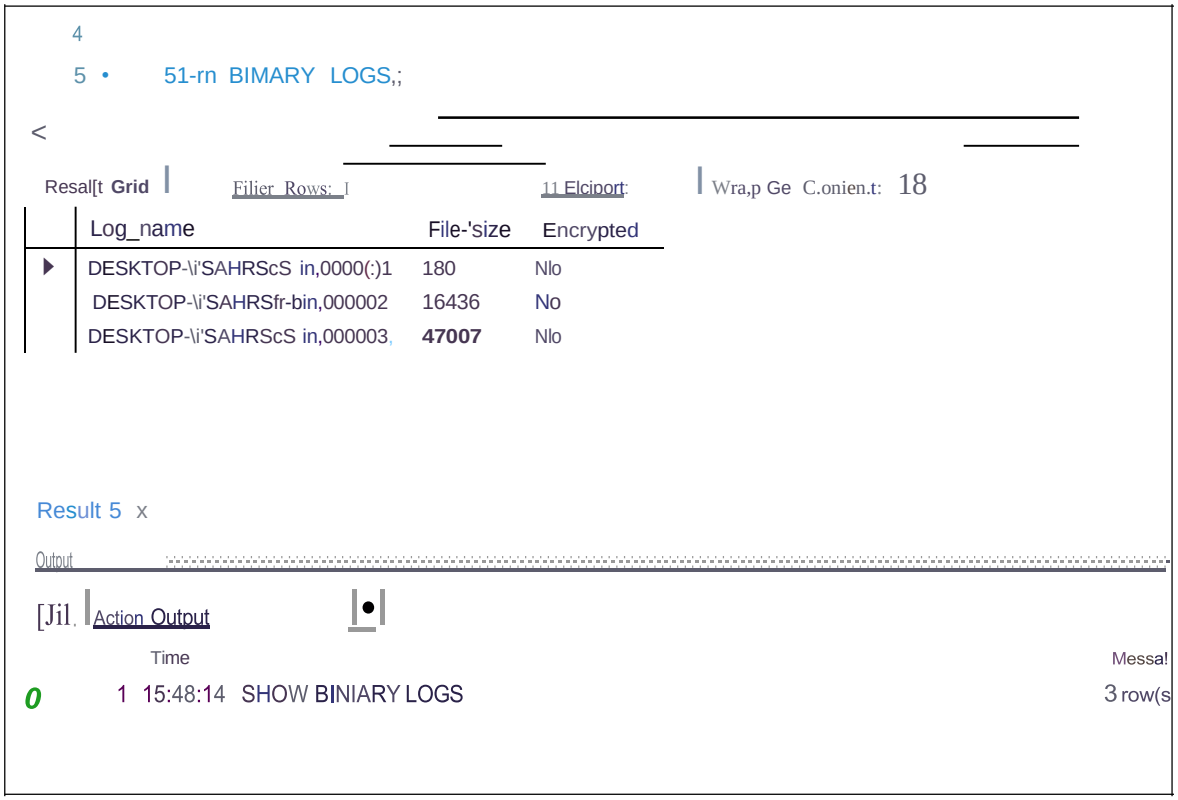
4
5 • 51-rn BIMARY LOGS,;
Resal[t Grid
Filier Rows: I
11 Elciport:
Wra,p Ge C.onien.t: 18
Log_name
File-'size
Encrypted
DESKTOP-\i'SAHRScS in,0000(:)1
180
Nlo
DESKTOP-\i'SAHRSfr-bin,000002
16436
No
DESKTOP-\i'SAHRScS in,000003,
47007
Nlo
Result 5 x
Output
[Jil
Action Output
Time
Messa!
0
1 15:48:14 SHOW BINIARY LOGS
3 row(s